

✓ Upload the ZIP in Colab

```
import zipfile
from google.colab import files
uploaded = files.upload()
```

```
with zipfile.ZipFile("archive (4).zip", 'r') as zip_ref:
    zip_ref.extractall("iris_data")
```



Choose Files archive (4).zip

- **archive (4).zip**(application/x-zip-compressed) - 1010 bytes, last modified: 7/1/2025 - 100% done
Saving archive (4).zip to archive (4) (1).zip

✓ Import Zip File

```
import zipfile
```

```
with zipfile.ZipFile("archive (4).zip", 'r') as zip_ref:
    zip_ref.extractall("iris_data")
```

✓ Extract the ZIP

```
import os
print(os.listdir("iris_data"))
```



['IRIS.csv']

✓ Load the CSV

```
import pandas as pd
df = pd.read_csv("iris_data/IRIS.csv")
print(df.head())
```



	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

✓ Show Details

```
print(df.isnull().sum())
# Check info
print(df.info())
```



```
sepal_length    0
sepal_width     0
petal_length    0
petal_width     0
species         0
dtype: int64
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   sepal_length    150 non-null   float64
1   sepal_width     150 non-null   float64
2   petal_length    150 non-null   float64
3   petal_width     150 non-null   float64
4   species         150 non-null   object
```

```
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
None
```

✓ Check for missing values

```
# Check missing values
print(df.isnull().sum())
```

```
↗ sepal_length    0
  sepal_width     0
  petal_length    0
  petal_width     0
  species         0
dtype: int64
```

✓ Inspect columns

```
print(df.columns)
```

```
↗ Index(['sepal_length', 'sepal_width', 'petal_length', 'petal_width',
        'species'],
        dtype='object')
```

✓ Drop unnecessary columns

```
if 'Id' in df.columns:
    df = df.drop('Id', axis=1)
```

✓ Encode the target label

```
from sklearn.preprocessing import LabelEncoder

le = LabelEncoder()
df['species'] = le.fit_transform(df['species'])

print(df['species'].value_counts())
```

```
↗ species
0    50
1    50
2    50
Name: count, dtype: int64
```

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Load the Iris dataset
from sklearn.datasets import load_iris
iris = load_iris()
df = pd.DataFrame(data=iris.data, columns=iris.feature_names)

# Create the correlation matrix
corr_matrix = df.corr()

# Plot the heatmap
plt.figure(figsize=(8,6))
sns.heatmap(corr_matrix, annot=True, cmap='Blues', linewidths=0.5)
plt.title("Feature Correlation Heatmap - Iris Dataset")
plt.show()
```



✓ 🌲 Random Forest Classifier

🌲 Random Forest Classifier

✓ Type: Ensemble method (many decision trees working together)

✓ How it works:

builds many decision trees on random subsets of data/features

averages their predictions (majority vote for classification)

each tree might overfit, but averaging them reduces overfitting

✓ Advantages:

works well with nonlinear and complex relationships

robust to outliers

handles missing data reasonably

works with minimal feature scaling

✓ Disadvantages:

more computationally heavy

harder to interpret

can overfit on small datasets without tuning.

✓ Train-test split

```
from sklearn.model_selection import train_test_split
from sklearn.datasets import load_iris

iris = load_iris()
X = pd.DataFrame(data=iris.data, columns=iris.feature_names)
y = pd.Series(data=iris.target, name='species')

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
```

✓ Train the Model Random Forest Classifier

```
from sklearn.ensemble import RandomForestClassifier

# initialize
clf = RandomForestClassifier(random_state=42)

# train
clf.fit(X_train, y_train)
```

↗

RandomForestClassifier ⓘ ?
 RandomForestClassifier(random_state=42)

✓ Evaluate the Model

```
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# predictions
y_pred = clf.predict(X_test)

# accuracy
print("Accuracy:", accuracy_score(y_test, y_pred))

# detailed report
print("\nClassification Report:\n", classification_report(y_test, y_pred))

# confusion matrix
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

↗ Accuracy: 0.9

Classification Report:		precision	recall	f1-score	support
0	1.00	1.00	1.00	10	
1	0.82	0.90	0.86	10	
2	0.89	0.80	0.84	10	
accuracy				0.90	30
macro avg		0.90	0.90	0.90	30
weighted avg		0.90	0.90	0.90	30

Confusion Matrix:

```
[[10  0  0]
 [ 0  9  1]
 [ 0  2  8]]
```

✓ + Logistic Regression

+ Logistic Regression

✓ Type: Linear model

✓ How it works:

fits a straight line (in high dimensions: a hyperplane)

predicts probabilities using a logistic/sigmoid function

uses coefficients for each feature

✓ Advantages:

easy to interpret (coefficients = feature impact)

fast to train

performs well if the data is linearly separable

very simple and robust

✔ Disadvantages:

poor performance on strongly nonlinear relationships

needs scaled features if large feature differences

underfits on complex patterns

✓ Train Logistic Regression

```
from sklearn.linear_model import LogisticRegression

# Logistic Regression with multinomial
logreg = LogisticRegression(multi_class='multinomial', solver='lbfgs', max_iter=200, random_state=42)

logreg.fit(X_train, y_train)

y_pred = logreg.predict(X_test)
```

⚠ /usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:1247: FutureWarning: 'multi_class' was deprecated in version 1
warnings.warn(

✓ Evaluate Logistic Regression

```
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

print("Logistic Regression Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

⚠ Logistic Regression Accuracy: 0.9666666666666667

Classification Report:					
	precision	recall	f1-score	support	
0	1.00	1.00	1.00	10	
1	1.00	0.90	0.95	10	
2	0.91	1.00	0.95	10	
accuracy			0.97	30	
macro avg	0.97	0.97	0.97	30	
weighted avg	0.97	0.97	0.97	30	

Confusion Matrix:

```
[[10  0  0]
 [ 0  9  1]
 [ 0  0 10]]
```

✓ Both Training and Testing accuracy

```
logreg = LogisticRegression(multi_class='multinomial', solver='lbfgs', max_iter=200, random_state=42)
logreg.fit(X_train, y_train)
```

⚠ /usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:1247: FutureWarning: 'multi_class' was deprecated in version 1
warnings.warn(

▼ LogisticRegression ⓘ ?
LogisticRegression(max_iter=200, multi_class='multinomial', random_state=42)

```
# predict on training data
y_train_pred = logreg.predict(X_train)
```

```
# accuracy on training set
train_acc = accuracy_score(y_train, y_train_pred)
print("Training Accuracy:", train_acc)

# accuracy on test set
test_acc = accuracy_score(y_test, y_pred)
print("Testing Accuracy:", test_acc)
```

↗ Training Accuracy: 0.975
Testing Accuracy: 0.9666666666666667

Decision Boundary Visualization for LogisticRegression And RandomForestClassifier

```
import matplotlib.pyplot as plt
import numpy as np
from matplotlib.colors import ListedColormap
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.datasets import load_iris
import pandas as pd

# Load the Iris dataset
iris = load_iris()
X = pd.DataFrame(data=iris.data, columns=iris.feature_names)
y = pd.Series(data=iris.target, name='species')

# choose 2 features
feature_x = 'petal length (cm)'
feature_y = 'petal width (cm)'

X_vis = X[[feature_x, feature_y]].values
y_vis = y

# train logistic regression on 2 features
logreg2d = LogisticRegression(multi_class='multinomial', solver='lbfgs', max_iter=200, random_state=42)
logreg2d.fit(X_vis, np.array(y_vis))

# train random forest on 2 features
rf2d = RandomForestClassifier(random_state=42)
rf2d.fit(X_vis, np.array(y_vis))

# meshgrid for visualization
x_min, x_max = X_vis[:, 0].min() - 1, X_vis[:, 0].max() + 1
y_min, y_max = X_vis[:, 1].min() - 1, X_vis[:, 1].max() + 1
xx, yy = np.meshgrid(
    np.arange(x_min, x_max, 0.02),
    np.arange(y_min, y_max, 0.02)
)

# predictions
Z_logreg = logreg2d.predict(np.c_[xx.ravel(), yy.ravel()]).reshape(xx.shape)
Z_rf = rf2d.predict(np.c_[xx.ravel(), yy.ravel()]).reshape(xx.shape)

# plot
plt.figure(figsize=(14,6))

# Logistic Regression plot
plt.subplot(1,2,1)
cmap_bg = ListedColormap(['#FFAAAA', '#AAFFAA', '#AAAAFF'])
cmap_pts = ListedColormap(['red', 'green', 'blue'])


plt.contourf(xx, yy, Z_logreg, alpha=0.4, cmap=cmap_bg)
plt.scatter(X_vis[:, 0], X_vis[:, 1], c=np.array(y_vis), cmap=cmap_pts, edgecolor='k', s=80)
plt.title("Logistic Regression Decision Boundary")
plt.xlabel(feature_x)
plt.ylabel(feature_y)

# Random Forest plot
plt.subplot(1,2,2)
```

```
plt.contourf(xx, yy, Z_rf, alpha=0.4, cmap=cmap_bg)
plt.scatter(X_vis[:, 0], X_vis[:, 1], c=np.array(y_vis), cmap=cmap_pts, edgecolor='k', s=80)
plt.title("Random Forest Decision Boundary")
plt.xlabel(feature_x)
plt.ylabel(feature_y)

plt.tight_layout()
plt.show()
```

→ /usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:1247: FutureWarning: 'multi_class' was deprecated in version 1
warnings.warn()



✓ Accuracy Comparison

```
# Assuming you already have these
logreg_accuracy = accuracy_score(y_test, logreg.predict(X_test)) # around 97%
rf_accuracy = accuracy_score(y_test, clf.predict(X_test)) # around 90%

# for clarity, print
print(f"Logistic Regression accuracy: {logreg_accuracy:.2f}")
print(f"Random Forest accuracy: {rf_accuracy:.2f}")
```

→ Logistic Regression accuracy: 0.97
Random Forest accuracy: 0.90

✓ Plot the bar graph

```
import matplotlib.pyplot as plt


models = ['Logistic Regression', 'Random Forest']
accuracies = [logreg_accuracy, rf_accuracy]

plt.figure(figsize=(8,6))
bars = plt.bar(models, accuracies, color=['skyblue', 'salmon'])

# add data labels
for bar in bars:
    yval = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2, yval + 0.01, f"{yval:.2f}", ha='center', fontsize=12)

plt.ylim(0, 1.1)
plt.ylabel("Accuracy")
plt.title("Model Accuracy Comparison")
```

```
plt.show()
```



✅ Quick Visualization Difference

Random Forest:

builds wiggly, irregular boundaries

great for complicated patterns

Logistic Regression:

draws straight lines

great for simple, clean, linearly separable data